

- [Dashboard](#)
- [Workbench](#)
- [Settings](#)

[Docs](#)

EB

4 AI contest, Click to rename

Get CodeRunCtrl + ↵

Model settings

Model

claude-3-opus-20240229

Max tokens to sample

[View all API options](#)

System Prompt

You are a senior chip design engineer who is an expert with Verilog HDL, especially designing AI accelerators for SoCs.

USER

I want you to help me join efabless'es 4 AI contest. The challenge aims to leverage generative AI to develop an open-source hardware accelerator designed explicitly for Keyword Spotting (KWS) applications on the Caravel System-on-Chip. Participants will utilize generative AI to optimize the KWS machine learning model and/or the audio features extractor (e.g., MFCC) to create an energy-efficient KWS accelerator that seamlessly integrates into the Caravel SoC environment. The discription is "In this contest, we invite you to utilize generative AI such as chatGPT, Gemini, Claude, Copilot, or similar tools to design a chipIgnite project targeted to an open-source hardware accelerator designed specifically for Keyword Spotting (KWS) applications.

Your design must be implemented using Efabless' chipIgnite that includes the Caravel SoC for rapid chip-level integration and the open-source design tools.

The objectives of this challenge go beyond creating designs; It's about raising awareness within the open source silicon design community about the wide range of possibilities for using Generative AI in chip design.

A successful project must provide all information necessary for other members of the community to reproduce your work. This includes detailed design documentation, the prompts used to generate the design, any required scripts or automation, and verification testbenches to demonstrate your design meets its intended functionality."

ASSISTANT

Certainly! I'd be happy to assist you in participating in the Efabless 4 AI contest. Let's break down the challenge and discuss the steps you can take to create an open-source hardware accelerator for Keyword Spotting (KWS) applications using generative AI.

1. Understand the requirements:

- Design an open-source hardware accelerator specifically for KWS applications.
- Utilize generative AI (e.g., ChatGPT, Gemini, Claude, Copilot) to optimize the KWS machine learning model and/or the audio features extractor (e.g., MFCC).
- Implement the design using Efabless' chipIgnite, which includes the Caravel SoC for rapid chip-level integration and open-source design tools.
- Provide detailed design documentation, prompts used, required scripts or automation, and verification testbenches.

2. Familiarize yourself with the tools and resources:

- Explore Efabless' chipIgnite platform and understand its capabilities and constraints.
- Familiarize yourself with the Caravel SoC and its integration process.
- Identify the open-source design tools provided by Efabless and learn how to use them effectively.

3. Design the KWS accelerator:

- Research existing KWS models and audio feature extractors suitable for hardware acceleration.
- Utilize generative AI to optimize the chosen KWS model and/or audio feature extractor for energy efficiency and hardware implementation.
- Develop the hardware architecture and design the accelerator using Verilog HDL or other supported hardware description languages.
- Integrate the accelerator with the Caravel SoC using the provided interfaces and protocols.

4. Document the design process:

- Keep detailed documentation of your design choices, optimizations, and trade-offs.
- Record the prompts and interactions with generative AI tools used during the design process.

- Provide clear instructions on how to reproduce your work, including any necessary scripts or automation.

5. Verify and validate the design:

- Create verification testbenches to ensure the functionality and performance of your KWS accelerator.
- Simulate and validate the design using the open-source tools provided by Efabless.
- Perform power, area, and timing analysis to assess the energy efficiency and resource utilization of your design.

6. Prepare the submission:

- Organize your design files, documentation, prompts, scripts, and testbenches into a clear and structured format.
- Ensure that your submission meets all the requirements and guidelines set by the contest organizers.
- Submit your project through the designated channels and within the specified deadlines.

Throughout the process, actively engage with the open-source silicon design community, seek feedback, and collaborate with others. The contest aims to raise awareness about the possibilities of using generative AI in chip design, so sharing your experiences and insights will be valuable to the community.

Remember to adhere to the contest rules, respect intellectual property rights, and provide proper attribution for any external resources or tools used in your design.

If you have any specific questions or need further guidance on any aspect of the contest or the design process, feel free to ask!

USER

I have already understood the requirements and I am already familiar with the tools and resources. Let's head straight to part 3 Design the KWS accelerator.

ASSISTANT

Great! Let's dive into designing the KWS accelerator. Here's a step-by-step approach to help you get started:

1. Choose a KWS model:

- Research and select a suitable KWS model for hardware acceleration, such as Deep Neural Networks (DNN), Convolutional Neural Networks (CNN), or Recurrent Neural Networks (RNN).
- Consider factors like model accuracy, complexity, and compatibility with hardware implementation.
- Utilize generative AI to optimize the chosen model architecture for energy efficiency and resource utilization.

2. Design the audio feature extractor:

- Identify an appropriate audio feature extractor, such as Mel-Frequency Cepstral Coefficients (MFCC), for preprocessing the audio input.
- Use generative AI to optimize the feature extractor algorithm for hardware implementation and energy efficiency.
- Develop the hardware architecture for the feature extractor, considering factors like parallelism, pipelining, and data reuse.

3. Develop the KWS accelerator architecture:

- Design the overall architecture of the KWS accelerator, including the datapath, control unit, and memory subsystem.
- Determine the data precision (e.g., fixed-point or floating-point) and bitwidth for the computations.
- Identify opportunities for parallelism, pipelining, and data reuse to improve performance and energy efficiency.
- Use generative AI to explore different architectural choices and optimize the design.

4. Implement the KWS accelerator in Verilog HDL:

- Translate the optimized KWS model and feature extractor into Verilog HDL modules.
- Implement the datapath, control unit, and memory subsystem using Verilog HDL.
- Ensure proper synchronization, data flow, and control signals between the different modules.
- Utilize generative AI to assist in generating Verilog code snippets and optimizing the HDL implementation.

5. Integrate the KWS accelerator with the Caravel SoC:

- Understand the interface and protocol requirements for integrating the accelerator with the Caravel SoC.
- Design the necessary interface logic and control signals to facilitate communication between the accelerator and the SoC.
- Ensure proper clock domain crossing and synchronization if the accelerator operates at a different frequency than the SoC.

6. Optimize the design for energy efficiency:

- Analyze the power consumption of the KWS accelerator using power estimation tools.
- Identify power-hungry modules or operations and explore techniques to reduce their power consumption.
- Apply clock gating, power gating, or dynamic voltage and frequency scaling (DVFS) to minimize power consumption during idle periods.
- Utilize generative AI to suggest power optimization techniques and generate power-efficient HDL code.

7. Verify the functionality and performance:

- Develop verification testbenches to validate the functionality of the KWS accelerator.
- Create test cases that cover different scenarios, including normal operation, edge cases, and error conditions.
- Simulate the design using the open-source tools provided by Efabless and ensure it meets the expected behavior.
- Measure the performance metrics, such as latency, throughput, and energy consumption, and compare them against the design goals.

Remember to document your design choices, optimizations, and the prompts used with generative AI throughout the process. This will be crucial for others to understand and reproduce your work.

If you have any specific questions or need further guidance on any aspect of the KWS accelerator design, feel free to ask!

USER

According to your knowledge select a suitable KWS model for hardware acceleration. Consider factors like model accuracy, complexity, and compatibility with hardware implementation.

ASSISTANT

Based on the requirements for hardware acceleration, model accuracy, complexity, and compatibility, I recommend using a Convolutional Neural Network (CNN) based model for the Keyword Spotting (KWS) accelerator. CNNs have shown excellent performance in speech recognition tasks, including KWS, and their architecture is well-suited for hardware implementation.

One specific CNN model that has been widely used for KWS is the CNN-KWS model, also known as the "Hello Edge" model. This model was introduced in the paper "Hello Edge: Keyword Spotting on Microcontrollers" by Zhang et al. (2017). The CNN-KWS model has several advantages:

1. Compact architecture: The model consists of a few convolutional layers followed by fully connected layers, making it relatively lightweight and suitable for resource-constrained hardware.
2. High accuracy: Despite its compact size, the CNN-KWS model achieves high accuracy in keyword spotting tasks, with reported accuracies of over 90% on popular KWS datasets like the Google Speech Commands dataset.
3. Compatibility with hardware: The convolutional and fully connected layers in the CNN-KWS model can be efficiently mapped to hardware resources like multiply-accumulate (MAC) units and memory buffers, enabling parallel and pipelined execution.
4. Energy efficiency: The compact size and hardware-friendly architecture of the CNN-KWS model make it energy-efficient, which is crucial for battery-powered devices and edge computing scenarios.

Here's a high-level overview of the CNN-KWS model architecture:

1. Input: Mel-frequency cepstral coefficients (MFCC) features extracted from the audio signal.
2. Convolutional layers: Two or three convolutional layers with small kernel sizes (e.g., 3x3) and a small number of filters (e.g., 32 or 64) to learn local patterns in the MFCC features.
3. Pooling layers: Max pooling layers to reduce the spatial dimensions and provide translation invariance.

4. Fully connected layers: One or two fully connected layers to learn high-level representations and perform classification.
5. Output layer: A softmax layer to produce the probability distribution over the keyword classes.

To optimize the CNN-KWS model for hardware implementation, you can consider the following techniques:

1. Quantization: Reduce the bitwidth of weights and activations to minimize storage and computation requirements. For example, using 8-bit fixed-point representation instead of 32-bit floating-point.
2. Pruning: Remove less significant weights or connections in the model to reduce complexity and memory footprint without significant accuracy loss.
3. Architectural optimizations: Explore variations of the CNN-KWS architecture, such as using depthwise separable convolutions or residual connections, to improve efficiency and accuracy.

You can utilize generative AI to assist in optimizing the CNN-KWS model architecture, suggesting quantization and pruning strategies, and generating optimized HDL code for hardware implementation.

Remember to train and validate the optimized CNN-KWS model on a suitable KWS dataset to ensure it meets the accuracy requirements before proceeding with hardware implementation.

If you have any further questions or need more details on any aspect of the KWS model selection and optimization, please let me know!

USER

Can you help me with a strategy to create a mel frequency cepstral coefficients (MFCCs) module?

Add Messages

5 of 5

Run prompt to see assistant response from Claude or...

Create a New Prompt